

BIOSTAT823 Final Project

December 12, 2022

Notebook by Caitlyn Nguyen

In this notebook, we explore different neural network architectures to predict for in-hospital mortality using the "In Hospital Mortality Prediction" data set, which is taken from <https://www.kaggle.com/datasets/saurabhshahane/in-hospital-mortality-prediction>. This dataset was derived from the MIMIC III dataset, and includes demographic characteristics (age, sex, ethnicity, weight, and height), vital signs (all continuous variables), comorbidities (all binary variables), and laboratory variables (all continuous variables).

The demographic characteristics and vital signs were taken once during the first 24 hours of ICU admission, laboratory variables were measured throughout the entire ICU stay, and comorbidities were determined via ICD-9 codes. The mean value was reported if there were multiple measurements of a variable recorded throughout the ICU stay. The primary outcome was in-hospital mortality, defined as the vital status of the patient at the time of discharge.

Training and validation epoch outputs are not shown as to limit the number of pages in the notebook.

1 Set-Up

1.1 Library Import

```
[ ]: # Library import
import os
import numpy as np
import math
import random
import matplotlib.pyplot as plt
import pandas as pd
import tensorflow as tf

from sklearn.impute import SimpleImputer
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import MinMaxScaler
from sklearn.metrics import mean_squared_error
from sklearn import metrics
from sklearn.utils import compute_class_weight
from tensorflow import keras
```

```

from tensorflow.keras import Sequential
from tensorflow.keras import layers
from tensorflow.keras.layers import Dense, Dropout
from tensorflow.keras import optimizers
from tensorflow.keras import regularizers
from tensorflow.keras import callbacks
from tensorflow.keras.callbacks import EarlyStopping
from sklearn.metrics import confusion_matrix
from sklearn.metrics import roc_curve
from sklearn.metrics import roc_auc_score
from sklearn.metrics import auc
from sklearn.metrics import precision_recall_curve
from sklearn.metrics import f1_score

```

1.2 Importing Data

The imported dataset had a size of 1177 x 51.

```

[ ]: # Set Seed
tf.random.set_seed(1)

```

```

[ ]: # All data
df = pd.read_csv("data01.csv")
print(df.shape)
df.head(5)

```

(1177, 51)

```

[ ]:
  group    ID  outcome  age  gendera      BMI  hypertensive  \
0      1  125047      0.0   72         1  37.588179          0
1      1  139812      0.0   75         2        NaN          0
2      1  109787      0.0   83         2  26.572634          0
3      1  130587      0.0   43         2  83.264629          0
4      1  138290      0.0   75         2  31.824842          1

  atrialfibrillation  CHD with no MI  diabetes  ...  Blood sodium  \
0                   0                0         1  ...   138.750000
1                   0                0         0  ...   138.888889
2                   0                0         0  ...   140.714286
3                   0                0         0  ...   138.500000
4                   0                0         0  ...   136.666667

  Blood calcium  Chloride  Anion gap  Magnesium ion    PH  Bicarbonate  \
0      7.463636  109.166667  13.166667      2.618182  7.230    21.166667
1      8.162500   98.444444  11.444444      1.887500  7.225    33.444444
2      8.266667  105.857143  10.000000      2.157143  7.268    30.571429

```

3	9.476923	92.071429	12.357143	1.942857	7.370	38.571429
4	8.733333	104.500000	15.166667	1.650000	7.250	22.000000

	Lactic acid	PCO2	EF
0	0.5	40.0	55
1	0.5	78.0	55
2	0.5	71.5	35
3	0.6	75.0	55
4	0.6	50.0	55

[5 rows x 51 columns]

2 Data Pre-processing

2.1 Data Scaling

We must first scale the continuous variables to be in range [0, 1] to standardize the image inputs for neural networks.

```
[ ]: # Scale continuous vars
noncont_cols = ["group", "ID", "outcome", "gendera", "hypertensive",
    ↳"atrialfibrillation", "CHD with no MI", "diabetes",
    ↳"deficiencyanemias", "depression", "Hyperlipemia", "Renal_
    ↳failure", "COPD"]
all_cols = list(df.columns.values)
cont_df = df.drop(noncont_cols, axis = 1)
scaled_df = (cont_df - cont_df.min()) / (cont_df.max() - cont_df.min())
clean_df = pd.concat((df[noncont_cols], scaled_df), 1)
print(clean_df.shape)
clean_df.head(5)
```

(1177, 51)

<ipython-input-27-a57081da2b61>:7: FutureWarning: In a future version of pandas all arguments of concat except for the argument 'objs' will be keyword-only

```
clean_df = pd.concat((df[noncont_cols], scaled_df), 1)
```

```
[ ]:   group      ID  outcome  gendera  hypertensive  atrialfibrillation  \
0      1  125047      0.0         1             0              0
1      1  139812      0.0         2             0              0
2      1  109787      0.0         2             0              0
3      1  130587      0.0         2             0              0
4      1  138290      0.0         2             1              0

      CHD with no MI  diabetes  deficiencyanemias  depression  ...  Blood sodium  \
0              0          1              1             0  ...      0.601029
```

1	0	0	1	0 ...	0.604495
2	0	0	1	0 ...	0.650050
3	0	0	0	0 ...	0.594790
4	0	0	1	0 ...	0.549037

	Blood calcium	Chloride	Anion gap	Magnesium ion	PH	Bicarbonate \
0	0.179679	0.683867	0.346185	0.455782	0.285714	0.238714
1	0.344118	0.430145	0.254886	0.182398	0.275510	0.591427
2	0.368627	0.605553	0.178313	0.283285	0.363265	0.508892
3	0.653394	0.279339	0.303270	0.203110	0.571429	0.738714
4	0.478431	0.573439	0.452209	0.093537	0.326531	0.262654

	Lactic acid	PCO2	EF
0	0.000000	0.266124	0.666667
1	0.000000	0.742016	0.666667
2	0.000000	0.660614	0.333333
3	0.012766	0.704446	0.666667
4	0.012766	0.391359	0.666667

[5 rows x 51 columns]

2.2 Missing Values

We then check for missing values. One observation had a missing value for outcome. That individual was removed. We then performed mean imputation for those with missing laboratory values.

```
[ ]: # Get number of NAs
clean_df.isna().sum()
```

```
[ ]: group          0
ID                  0
outcome            1
gendera            0
hypertensive        0
atrialfibrillation  0
CHD with no MI      0
diabetes             0
deficiencyanemias   0
depression           0
Hyperlipemia         0
Renal failure         0
COPD                 0
age                  0
BMI                  215
heart rate            13
Systolic blood pressure 16
```

Diastolic blood pressure	16
Respiratory rate	13
temperature	19
SP O2	13
Urine output	36
hematocrit	0
RBC	0
MCH	0
MCHC	0
MCV	0
RDW	0
Leucocyte	0
Platelets	0
Neutrophils	144
Basophils	259
Lymphocyte	145
PT	20
INR	20
NT-proBNP	0
Creatine kinase	165
Creatinine	0
Urea nitrogen	0
glucose	18
Blood potassium	0
Blood sodium	0
Blood calcium	1
Chloride	0
Anion gap	0
Magnesium ion	0
PH	292
Bicarbonate	0
Lactic acid	229
PCO2	294
EF	0
dtype: int64	

```
[ ]: # Remove where outcome is unknown
clean_df2 = clean_df[~clean_df['outcome'].isna()]
print(clean_df2.shape)
clean_df2.head(5)
```

(1176, 51)

```
[ ]:   group    ID  outcome  gendera  hypertensive  atrialfibrillation  \
0      1  125047     0.0         1             0              0
1      1  139812     0.0         2             0              0
2      1  109787     0.0         2             0              0
```

3	1	130587	0.0	2	0	0
4	1	138290	0.0	2	1	0

	CHD with no MI	diabetes	deficiencyanemias	depression	...	Blood sodium	\
0	0	1	1	0	...	0.601029	
1	0	0	1	0	...	0.604495	
2	0	0	1	0	...	0.650050	
3	0	0	0	0	...	0.594790	
4	0	0	1	0	...	0.549037	

	Blood calcium	Chloride	Anion gap	Magnesium ion	PH	Bicarbonate	\
0	0.179679	0.683867	0.346185	0.455782	0.285714	0.238714	
1	0.344118	0.430145	0.254886	0.182398	0.275510	0.591427	
2	0.368627	0.605553	0.178313	0.283285	0.363265	0.508892	
3	0.653394	0.279339	0.303270	0.203110	0.571429	0.738714	
4	0.478431	0.573439	0.452209	0.093537	0.326531	0.262654	

	Lactic acid	PCO2	EF
0	0.000000	0.266124	0.666667
1	0.000000	0.742016	0.666667
2	0.000000	0.660614	0.333333
3	0.012766	0.704446	0.666667
4	0.012766	0.391359	0.666667

[5 rows x 51 columns]

```
[ ]: # Imput with mean
cont_df = clean_df2.drop(noncont_cols, axis = 1)
imputed_df = cont_df.fillna(cont_df.mean())
clean_df3 = pd.concat((clean_df2[noncont_cols], imputed_df), 1)
clean_df3.head(5)
```

<ipython-input-183-601ef2ee5cec>:4: FutureWarning: In a future version of pandas all arguments of concat except for the argument 'objs' will be keyword-only

```
clean_df3 = pd.concat((clean_df2[noncont_cols], imputed_df), 1)
```

```
[ ]: group      ID outcome  gendera  hypertensive  atrialfibrillation  \
0      1  125047      0.0        1            0            0
1      1  139812      0.0        2            0            0
2      1  109787      0.0        2            0            0
3      1  130587      0.0        2            0            0
4      1  138290      0.0        2            1            0

      CHD with no MI  diabetes  deficiencyanemias  depression  ...  Blood sodium  \
0              0        1              1            0  ...    0.601029
1              0        0              1            0  ...    0.604495
2              0        0              1            0  ...    0.650050
```

3	0	0	0	0	...	0.594790
4	0	0	1	0	...	0.549037

	Blood calcium	Chloride	Anion gap	Magnesium ion	PH	Bicarbonate \
0	0.179679	0.683867	0.346185	0.455782	0.285714	0.238714
1	0.344118	0.430145	0.254886	0.182398	0.275510	0.591427
2	0.368627	0.605553	0.178313	0.283285	0.363265	0.508892
3	0.653394	0.279339	0.303270	0.203110	0.571429	0.738714
4	0.478431	0.573439	0.452209	0.093537	0.326531	0.262654

	Lactic acid	PCO2	EF
0	0.000000	0.266124	0.666667
1	0.000000	0.742016	0.666667
2	0.000000	0.660614	0.333333
3	0.012766	0.704446	0.666667
4	0.012766	0.391359	0.666667

[5 rows x 51 columns]

2.3 Data types

We then looked at the data types of the variables. The outcome variable was shown to be of float64. We converted it to int64. We also saw that gender was coded as 1 or 2. We converted gender to be coded as 0 or 1.

```
[ ]: # Get info
print(clean_df3.info(verbose = True))
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 1176 entries, 0 to 1176
Data columns (total 51 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   group                                1176 non-null   int64
1   ID                                    1176 non-null   int64
2   outcome                              1176 non-null   float64
3   gendera                              1176 non-null   int64
4   hypertensive                         1176 non-null   int64
5   atrialfibrillation                   1176 non-null   int64
6   CHD with no MI                       1176 non-null   int64
7   diabetes                             1176 non-null   int64
8   deficiencyanemias                    1176 non-null   int64
9   depression                           1176 non-null   int64
10  Hyperlipemia                          1176 non-null   int64
11  Renal failure                         1176 non-null   int64
12  COPD                                  1176 non-null   int64
13  age                                   1176 non-null   float64
```

14	BMI	1176 non-null	float64
15	heart rate	1176 non-null	float64
16	Systolic blood pressure	1176 non-null	float64
17	Diastolic blood pressure	1176 non-null	float64
18	Respiratory rate	1176 non-null	float64
19	temperature	1176 non-null	float64
20	SP O2	1176 non-null	float64
21	Urine output	1176 non-null	float64
22	hematocrit	1176 non-null	float64
23	RBC	1176 non-null	float64
24	MCH	1176 non-null	float64
25	MCHC	1176 non-null	float64
26	MCV	1176 non-null	float64
27	RDW	1176 non-null	float64
28	Leucocyte	1176 non-null	float64
29	Platelets	1176 non-null	float64
30	Neutrophils	1176 non-null	float64
31	Basophils	1176 non-null	float64
32	Lymphocyte	1176 non-null	float64
33	PT	1176 non-null	float64
34	INR	1176 non-null	float64
35	NT-proBNP	1176 non-null	float64
36	Creatine kinase	1176 non-null	float64
37	Creatinine	1176 non-null	float64
38	Urea nitrogen	1176 non-null	float64
39	glucose	1176 non-null	float64
40	Blood potassium	1176 non-null	float64
41	Blood sodium	1176 non-null	float64
42	Blood calcium	1176 non-null	float64
43	Chloride	1176 non-null	float64
44	Anion gap	1176 non-null	float64
45	Magnesium ion	1176 non-null	float64
46	PH	1176 non-null	float64
47	Bicarbonate	1176 non-null	float64
48	Lactic acid	1176 non-null	float64
49	PCO2	1176 non-null	float64
50	EF	1176 non-null	float64

dtypes: float64(39), int64(12)
memory usage: 477.8 KB
None

```
[ ]: # Change dtype of outcome
clean_df3.outcome = clean_df3.outcome.astype(int)
print(clean_df3.info(verbose = True))
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 1176 entries, 0 to 1176
Data columns (total 51 columns):
```


#	Column	Non-Null Count	Dtype
----	-----	-----	-----
0	group	1176 non-null	int64
1	ID	1176 non-null	int64
2	outcome	1176 non-null	int64
3	gendera	1176 non-null	int64
4	hypertensive	1176 non-null	int64
5	atrialfibrillation	1176 non-null	int64
6	CHD with no MI	1176 non-null	int64
7	diabetes	1176 non-null	int64
8	deficiencyanemias	1176 non-null	int64
9	depression	1176 non-null	int64
10	Hyperlipemia	1176 non-null	int64
11	Renal failure	1176 non-null	int64
12	COPD	1176 non-null	int64
13	age	1176 non-null	float64
14	BMI	1176 non-null	float64
15	heart rate	1176 non-null	float64
16	Systolic blood pressure	1176 non-null	float64
17	Diastolic blood pressure	1176 non-null	float64
18	Respiratory rate	1176 non-null	float64
19	temperature	1176 non-null	float64
20	SP O2	1176 non-null	float64
21	Urine output	1176 non-null	float64
22	hematocrit	1176 non-null	float64
23	RBC	1176 non-null	float64
24	MCH	1176 non-null	float64
25	MCHC	1176 non-null	float64
26	MCV	1176 non-null	float64
27	RDW	1176 non-null	float64
28	Leucocyte	1176 non-null	float64
29	Platelets	1176 non-null	float64
30	Neutrophils	1176 non-null	float64
31	Basophils	1176 non-null	float64
32	Lymphocyte	1176 non-null	float64
33	PT	1176 non-null	float64
34	INR	1176 non-null	float64
35	NT-proBNP	1176 non-null	float64
36	Creatine kinase	1176 non-null	float64
37	Creatinine	1176 non-null	float64
38	Urea nitrogen	1176 non-null	float64
39	glucose	1176 non-null	float64
40	Blood potassium	1176 non-null	float64
41	Blood sodium	1176 non-null	float64
42	Blood calcium	1176 non-null	float64
43	Chloride	1176 non-null	float64
44	Anion gap	1176 non-null	float64
45	Magnesium ion	1176 non-null	float64

```

46 PH 1176 non-null float64
47 Bicarbonate 1176 non-null float64
48 Lactic acid 1176 non-null float64
49 PCO2 1176 non-null float64
50 EF 1176 non-null float64
dtypes: float64(38), int64(13)
memory usage: 477.8 KB
None

```

```

[ ]: # Change gender to be binary, 0 or 1
clean_df4 = clean_df3.copy()
clean_df4['gendera'] = clean_df4['gendera'].replace([1,2],[0, 1])
clean_df4.gendera.head(5)

```

```

[ ]: 0    1017
     1     159
     Name: outcome, dtype: int64

```

3 Dataset Splitting

The dataset was split to be in a 80:10:10 ratio (training, validation, testing). The split data was then converted to tensors, with the predictor tensors of the training, validation, and testing datasets being of respective size (940, 48), (118, 48), and (118, 48). Class imbalance was checked and in-hospital mortality was shown to have a prevalence of 0.135. To address this, class weights were calculated to give more weight to those observations with an observed in-hospital mortality.

```

[ ]: # Split train, val, and test in 80:10:10
x = clean_df4.drop(["group", "ID", "outcome"], axis = 1).copy()
y = clean_df4['outcome']
X_train, X_rem, y_train, y_rem = train_test_split(x,y, train_size=0.8)
X_valid, X_test, y_valid, y_test = train_test_split(X_rem,y_rem, test_size=0.5)

print(X_train.shape), print(y_train.shape)
print(X_valid.shape), print(y_valid.shape)
print(X_test.shape), print(y_test.shape)

```

```

(940, 48)
(940,)
(118, 48)
(118,)
(118, 48)
(118,)

```

```

[ ]: (None, None)

```

```
[ ]: # Convert to tensors
train_x = tf.stack(X_train)
train_y = tf.stack(y_train)
val_x = tf.stack(X_valid)
val_y = tf.stack(y_valid)
test_x = tf.stack(X_test)
test_y = tf.stack(y_test)
test_y
```

```
[ ]: <tf.Tensor: shape=(118,), dtype=int64, numpy=
array([0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 1, 0, 0, 0, 0, 0,
       1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0,
       0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 1, 0, 0, 1, 0, 0, 0,
       1, 0, 1, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 1, 1, 0, 0, 0, 0, 0, 0,
       0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1,
       0, 0, 1, 0, 0, 0, 0, 0])>
```

```
[ ]: # Examine class imbalance
clean_df4['outcome'].value_counts()
```

```
[ ]: 0    1017
     1     159
     Name: outcome, dtype: int64
```

```
[ ]: class_weights = compute_class_weight(class_weight = "balanced",
                                         classes = np.unique(y_train),
                                         y = y_train
                                         )
class_weights = dict(zip(np.unique(y_train), class_weights))
class_weights
```

```
[ ]: {0: 0.5752753977968176, 1: 3.821138211382114}
```

4 First Model

4.1 Architecture

```
[ ]: model = Sequential([
    tf.keras.Input(48),
    tf.keras.layers.Flatten(),
    tf.keras.layers.Dense(128, activation=tf.nn.relu),
    tf.keras.layers.Dense(1, activation=tf.nn.sigmoid)
])
model.summary()
```

Model: "sequential_33"

Layer (type)	Output Shape	Param #
flatten_34 (Flatten)	(None, 48)	0
dense_68 (Dense)	(None, 128)	6272
dense_69 (Dense)	(None, 1)	129

Total params: 6,401

Trainable params: 6,401

Non-trainable params: 0

The model has an input layer, flatten layer, and two dense layers. The first input layers takes in the shape of (48) as that is the shape of the number of predictors. Inputs are flattened and passed to a dense layer with 128 hidden units and a ReLU activation function. The final output layer has 1 hidden unit with activation function sigmoid to return a predicted probability for in-hospital mortality from 0 to 1.

```
[ ]: # Define path to store callback
cp_path = "model1.ckpt"
cp_dir = os.path.dirname(cp_path)

# Callback that saves the models weights at the best validation accuracy
cp_weights = tf.keras.callbacks.ModelCheckpoint(filepath=cp_path,
                                                save_weights_only=False,
                                                monitor='val_accuracy',
                                                mode='max',
                                                save_best_only=True)

# Compile model
model.compile(optimizer = optimizers.Adam(0.001), loss = 'binary_crossentropy',
              metrics = ['accuracy'])
```

The model was compiled with an Adam optimizer with learning rate = 0.001, loss calculated through the binary cross entropy, and the evaluated metric to be accuracy.

4.2 Training and Validation

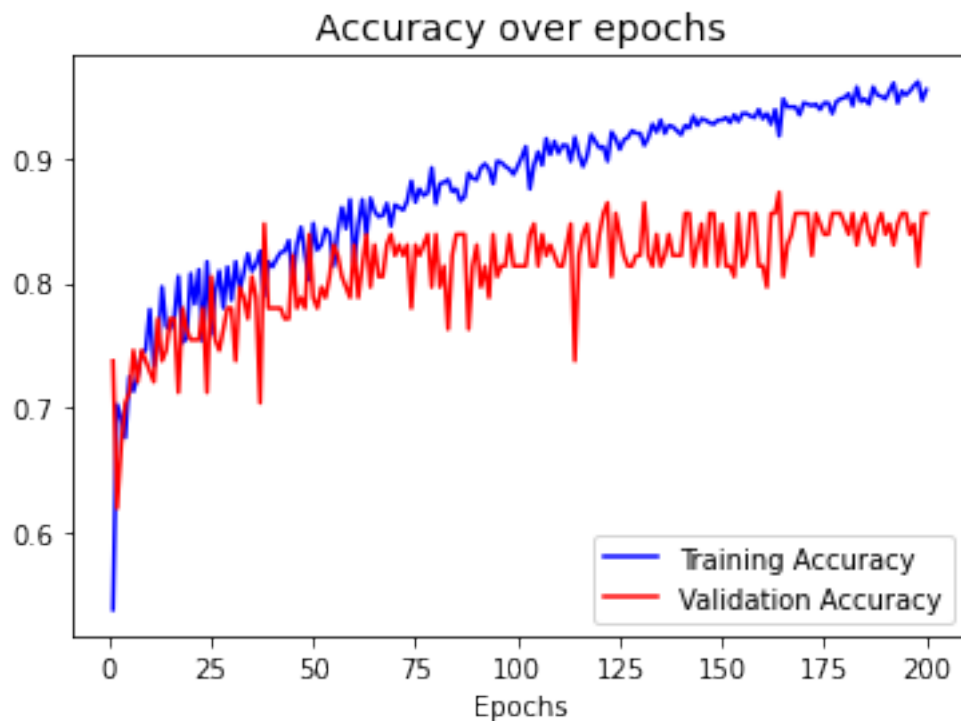
```
[ ]: history = model.fit(train_x, train_y, epochs = 200, batch_size=64,
                        validation_data=(val_x, val_y), callbacks=[cp_weights],
                        class_weight=class_weights, verbose=1)
```

4.3 Training Evaluation

```
[ ]: # Obtain training and validation metrics
train_acc = history.history['accuracy']
val_acc = history.history['val_accuracy']
train_loss = history.history['loss']
val_loss = history.history['val_loss']
epochs = range(1, len(train_acc) + 1)

# Accuracy over epochs
plt.plot(epochs, train_acc, color='blue', label='Training Accuracy')
plt.plot(epochs, val_acc, color='red', label='Validation Accuracy')
plt.title("Accuracy over epochs", fontsize=14)
plt.xlabel('Epochs')
plt.legend(loc='lower right')
plt.show()

# Print best validation accuracy and epoch
print("Best validation accuracy was %5.3f at epoch %2.f" % (np.max(val_acc), np.
    ↳argmax(val_acc)+1))
```

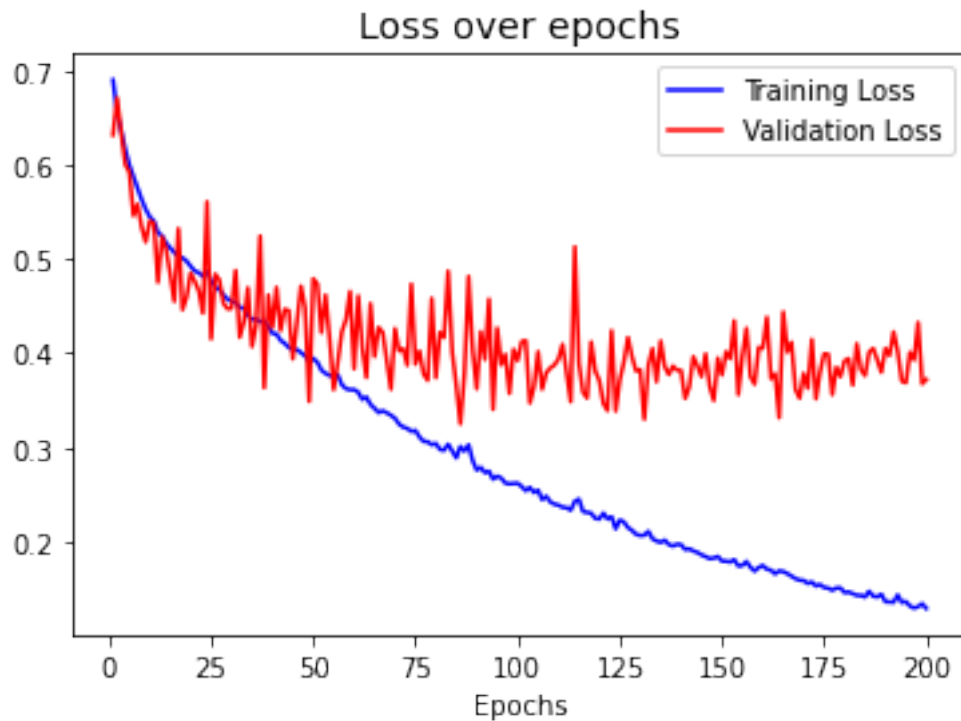


Best validation accuracy was 0.873 at epoch 164

The best validation accuracy was observed to be 0.873 at epoch 164. The Accuracy vs Epoch plot

shows that there is overfitting observed beginning at epoch 70 because following that epoch, the validation accuracy plateaus while the training accuracy continues to increase.

```
[ ]: # Loss over epochs
plt.plot(epochs, train_loss, color='blue', label='Training Loss')
plt.plot(epochs, val_loss, color='red', label='Validation Loss')
plt.title("Loss over epochs", fontsize=14)
plt.xlabel('Epochs')
plt.legend(loc='upper right')
plt.show()
```



Overfitting is observed with the Loss vs Epoch plot. At around epoch 20, the validation loss begins to plateau while the training loss continues to decrease.

5 Second Model

5.1 Architecture

```
[ ]: model2 = Sequential([
    tf.keras.Input(48),
    tf.keras.layers.Dropout(0.05),
    tf.keras.layers.Flatten(),
```

```

        tf.keras.layers.Dense(128, activation=tf.nn.relu),
        tf.keras.layers.Dense(1, activation=tf.nn.sigmoid)
    ]
)
model2.summary()

```

Model: "sequential_35"

Layer (type)	Output Shape	Param #
dropout_23 (Dropout)	(None, 48)	0
flatten_36 (Flatten)	(None, 48)	0
dense_72 (Dense)	(None, 128)	6272
dense_73 (Dense)	(None, 1)	129

=====
 Total params: 6,401
 Trainable params: 6,401
 Non-trainable params: 0
 =====

As the model was overfitting, adding a dropout layer would help. A dropout layer was added before flattening to drop out observations with a rate of 0.05.

```

[ ]: # Define path to store callback
cp_path = "model2.ckpt"
cp_dir = os.path.dirname(cp_path)

# Callback that saves the models weights at the best validation accuracy
cp_weights = tf.keras.callbacks.ModelCheckpoint(filepath=cp_path,
                                                save_weights_only=False,
                                                monitor='val_accuracy',
                                                mode='max',
                                                save_best_only=True)

# Compile model
model2.compile(optimizer = optimizers.Adam(0.0001), loss =_
               ↪ 'binary_crossentropy', metrics = ['accuracy'])

```

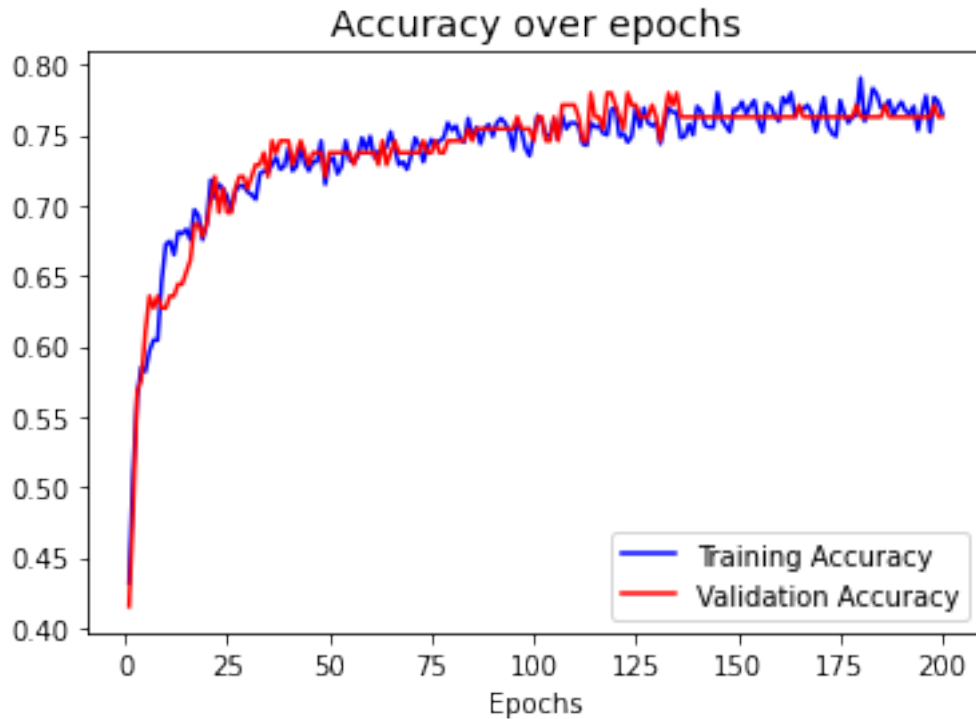
The learning rate was decreased from 0.001 to 0.0001 to slow the learning of the model, since within 50 epochs the model already began to overfit.

5.2 Training and Validation

```
[ ]: history2 = model2.fit(train_x, train_y, epochs = 200, batch_size=64,  
    ↪ validation_data=(val_x, val_y), callbacks=[cp_weights],  
    ↪ class_weight=class_weights, verbose=1)
```

5.3 Training Evaluation

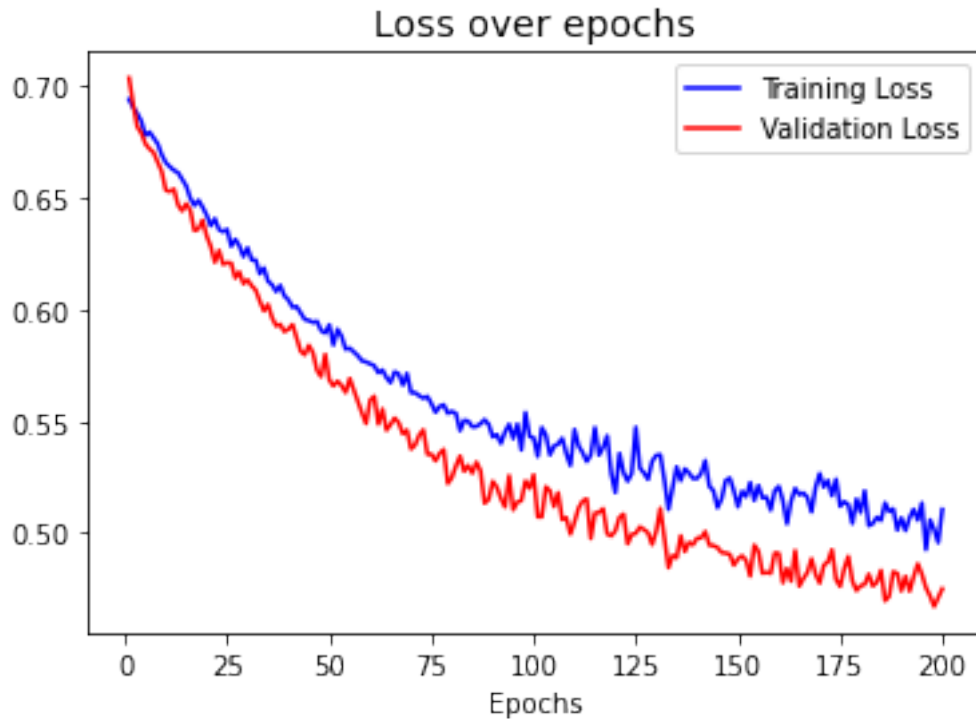
```
[ ]: # Obtain training and validation metrics  
train_acc = history2.history['accuracy']  
val_acc = history2.history['val_accuracy']  
train_loss = history2.history['loss']  
val_loss = history2.history['val_loss']  
epochs = range(1, len(train_acc) + 1)  
  
# Accuracy over epochs  
plt.plot(epochs, train_acc, color='blue', label='Training Accuracy')  
plt.plot(epochs, val_acc, color='red', label='Validation Accuracy')  
plt.title("Accuracy over epochs", fontsize=14)  
plt.xlabel('Epochs')  
plt.legend(loc='lower right')  
plt.show()  
  
# Print best validation accuracy and epoch  
print("Best validation accuracy was %5.3f at epoch %2.f" % (np.max(val_acc), np.  
    ↪ argmax(val_acc) + 1))
```

Best validation accuracy was 0.780 at epoch 114

The best accuracy for the validation set was 0.780 at epoch 114. This validation accuracy is worse than the first validation accuracy of 0.873 which occurred at epoch 164. It appears that the training and validation accuracy have plateaued around 0.750.

```
[ ]: # Loss over epochs
plt.plot(epochs, train_loss, color='blue', label='Training Loss')
plt.plot(epochs, val_loss, color='red', label='Validation Loss')
plt.title("Loss over epochs", fontsize=14)
plt.xlabel('Epochs')
plt.legend(loc='upper right')
plt.show()
```



Slight underfitting is observed. Both losses are still decreasing (although validation is decreasing at a slower rate), indicating that there is underfitting and more training can be done to reach the optimum solution.

6 Third Model

6.1 Architecture

```
[ ]: model3 = Sequential([
    tf.keras.Input(48),
    tf.keras.layers.Dropout(0.1),
    tf.keras.layers.Flatten(),
    tf.keras.layers.Dense(128, activation=tf.nn.relu),
    tf.keras.layers.Dense(64, activation=tf.nn.relu),
    tf.keras.layers.Dense(32, activation=tf.nn.relu),
    tf.keras.layers.Dense(1, activation=tf.nn.sigmoid)
])
model3.summary()
```

Model: "sequential_43"

Layer (type)	Output Shape	Param #
dropout_31 (Dropout)	(None, 48)	0
flatten_44 (Flatten)	(None, 48)	0
dense_100 (Dense)	(None, 128)	6272
dense_101 (Dense)	(None, 64)	8256
dense_102 (Dense)	(None, 32)	2080
dense_103 (Dense)	(None, 1)	33

```

=====
Total params: 16,641
Trainable params: 16,641
Non-trainable params: 0
-----

```

Due to some slight potential for underfitting, a third and fourth dense layer were added, with 64 and 32 hidden units. This was done to increase complexity of the model to address the slight potential for underfitting. To slightly offset the added complexity of the model, the dropout rate was increased from 0.05 to 0.1.

```

[ ]: # Define path to store callback
cp_path = "model3.ckpt"
cp_dir = os.path.dirname(cp_path)

# Callback that saves the models weights at the best validation accuracy
cp_weights = tf.keras.callbacks.ModelCheckpoint(filepath=cp_path,
                                                save_weights_only=False,
                                                monitor='val_accuracy',
                                                mode='max',
                                                save_best_only=True)

# Compile model
model3.compile(optimizer = optimizers.Adam(0.0001), loss =
↳ 'binary_crossentropy', metrics = ['accuracy'])

```

6.2 Training and Validation

```

[ ]: history3 = model3.fit(train_x, train_y, epochs = 200, batch_size=64,
↳ validation_data=(val_x, val_y), callbacks=[cp_weights],
↳ class_weight=class_weights, verbose=1)

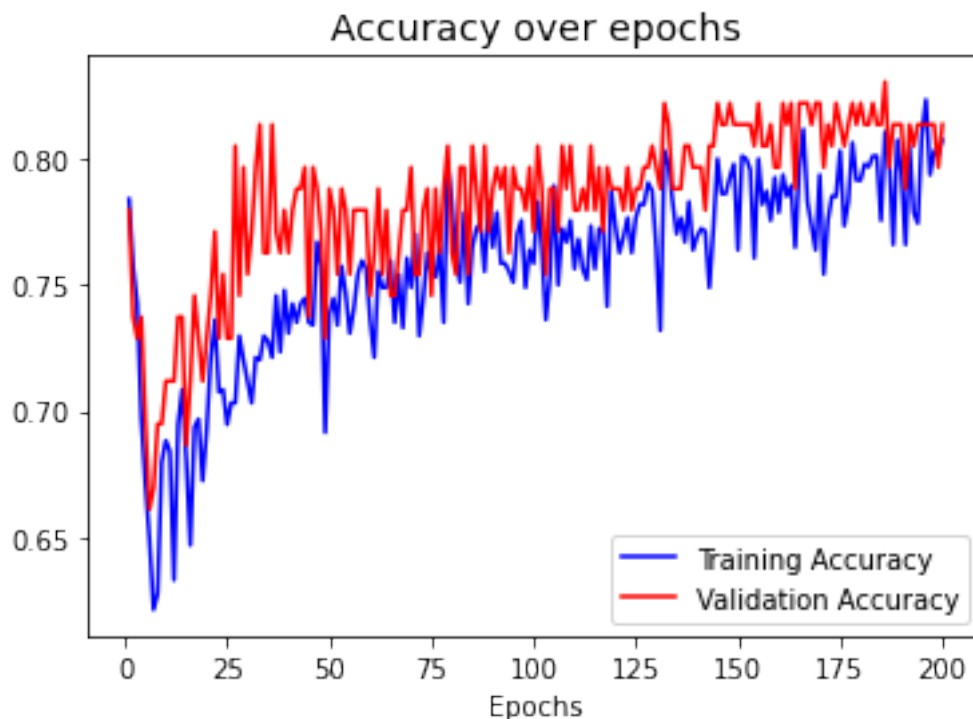
```

6.3 Training Evaluation

```
[ ]: # Obtain training and validation metrics
train_acc = history3.history['accuracy']
val_acc = history3.history['val_accuracy']
train_loss = history3.history['loss']
val_loss = history3.history['val_loss']
epochs = range(1, len(train_acc) + 1)

# Accuracy over epochs
plt.plot(epochs, train_acc, color='blue', label='Training Accuracy')
plt.plot(epochs, val_acc, color='red', label='Validation Accuracy')
plt.title("Accuracy over epochs", fontsize=14)
plt.xlabel('Epochs')
plt.legend(loc='lower right')
plt.show()

# Print best validation accuracy and epoch
print("Best validation accuracy was %5.3f at epoch %2.f" % (np.max(val_acc), np.
    ↳argmax(val_acc)+1))
```

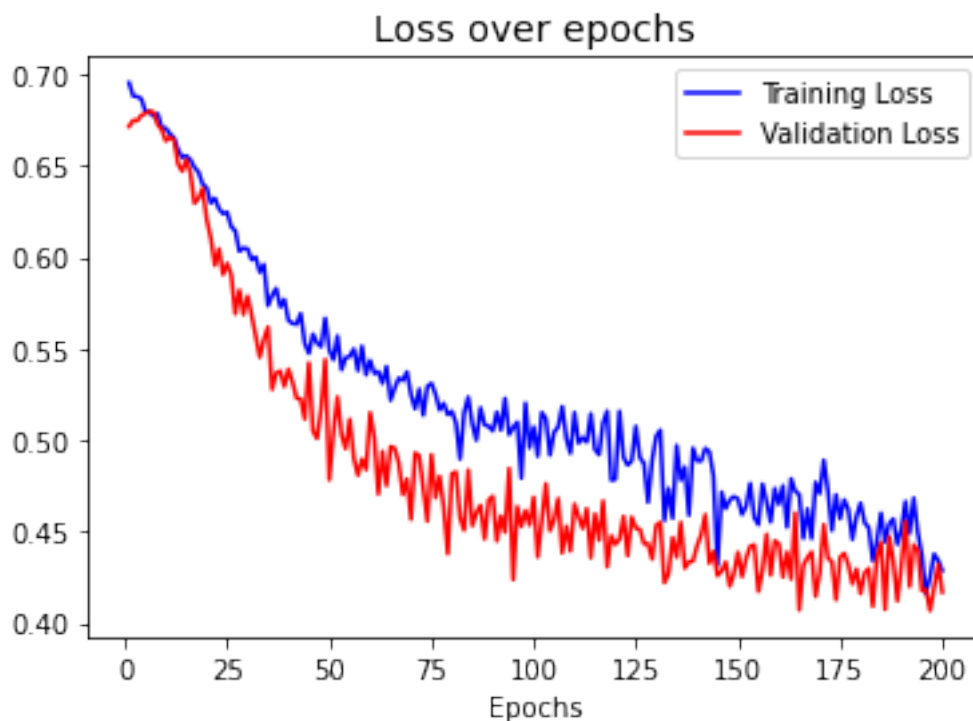


Best validation accuracy was 0.831 at epoch 186

The best validation accuracy was 0.831 at epoch 186, which was less than the validation accuracy

of the first model architecture (0.873 at epoch 164) but greater than the validation accuracy of the second model architecture (0.780 at epoch 114). The accuracies began high, most likely due to chance in the mini-batches, and then decreased within the first 5 epochs. Afterwards, the accuracies began to increase. The validation accuracy and training accuracy appear to be close together. More training time could potentially be done as the accuracies somewhat seem to have potential to continue increasing.

```
[ ]: # Loss over epochs
plt.plot(epochs, train_loss, color='blue', label='Training Loss')
plt.plot(epochs, val_loss, color='red', label='Validation Loss')
plt.title("Loss over epochs", fontsize=14)
plt.xlabel('Epochs')
plt.legend(loc='upper right')
plt.show()
```



The validation loss was consistently below the training loss. The loss in the validation seems to be around 0.40 by the end of the training, which matches that of the training loss. There could be potential for more training if there were more epochs.

7 Final Model Evaluation

The best model architecture appeared to be the third model architecture. This model had the second greatest validation accuracy (0.831) as compared to the other model architectures (0.873 and

0.780, respectively). The first model architecture observed overfitting, as the validation accuracy begin to deviate from the training accuracy at around epoch 70, then the validation accuracy plateaued afterwards. The second architecture had the slight potential to train more as it was slightly underfitting, where the losses were still trending downwards. The third architecture similar accuracy and loss between the training and testing dataset, and had a lower validation loss than the second model. I believe that the third model is the best model as the first architecture had much overfitting, and the second model architecture had slight underfitting, whereas the third model improved accuracy from the second model while also retaining close accuracy and loss between the training and validation set.

7.1 Obtain Best Model

```
[ ]: # Get weights of models from epoch with best val accuracy
model.load_weights("model1.ckpt")
model2.load_weights("model2.ckpt")
model3.load_weights("model3.ckpt")

# Save models to local drive in case Google Colab crashes
!mkdir -p saved_model
model3.save('saved_model/model3')
model2.save('saved_model/model1')
model.save('saved_model/model2')
```

7.2 Accuracy

```
[ ]: # Get accuracy on testing set
loss, accuracy = model3.evaluate(test_x, test_y)
print('Accuracy on test dataset: %5.3f' % (accuracy))
```

```
4/4 [=====] - 0s 9ms/step - loss: 0.4281 - accuracy:
0.8136
Accuracy on test dataset: 0.814
```

7.3 Confusion Matrix

```
[ ]: # Get predictions
test_pred = model3.predict(test_x)
```

```
4/4 [=====] - 0s 3ms/step
```

```
[ ]: # Check predictions
test_pred[0:5]
```

```
[ ]: array([[0.1500526 ],
          [0.18358165],
          [0.03791126],
          [0.14376739],
          [0.8128861 ]], dtype=float32)
```

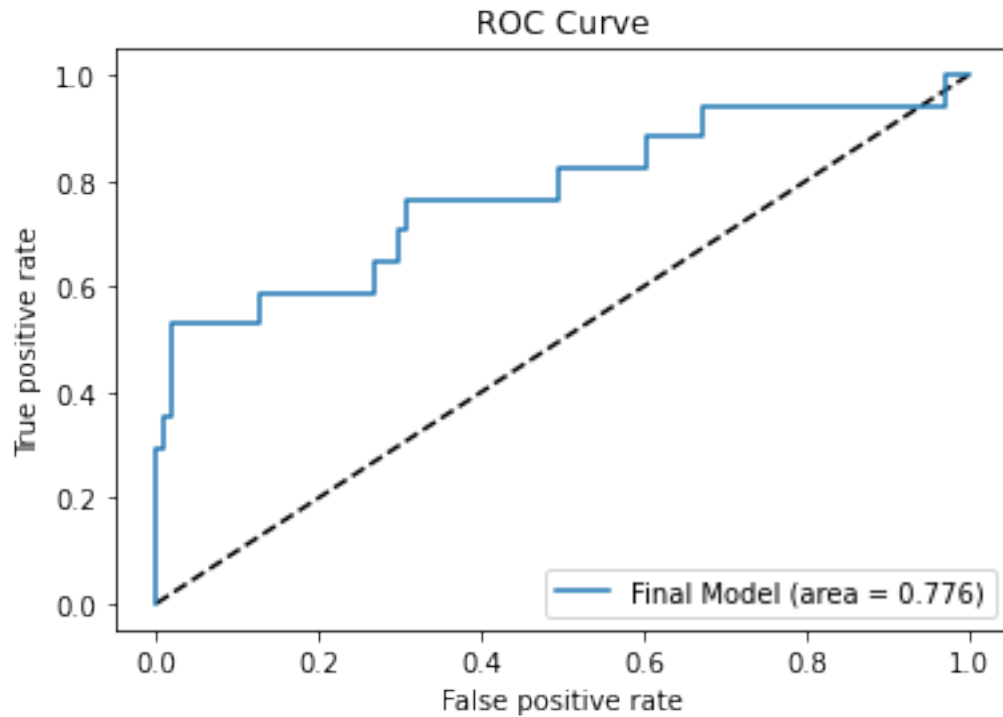
```
[ ]: # Get confusion matrix
conf_matrix = tf.math.confusion_matrix(labels=test_y, predictions=np.
    ↪rint(test_pred))
conf_matrix
```

```
[ ]: <tf.Tensor: shape=(2, 2), dtype=int32, numpy=
    array([[86, 15],
           [ 7, 10]], dtype=int32)>
```

7.4 ROC

```
[ ]: # Get ROC curve variables
fpr, tpr, thresh = roc_curve(test_y, test_pred)
model_auc = auc(fpr, tpr)
```

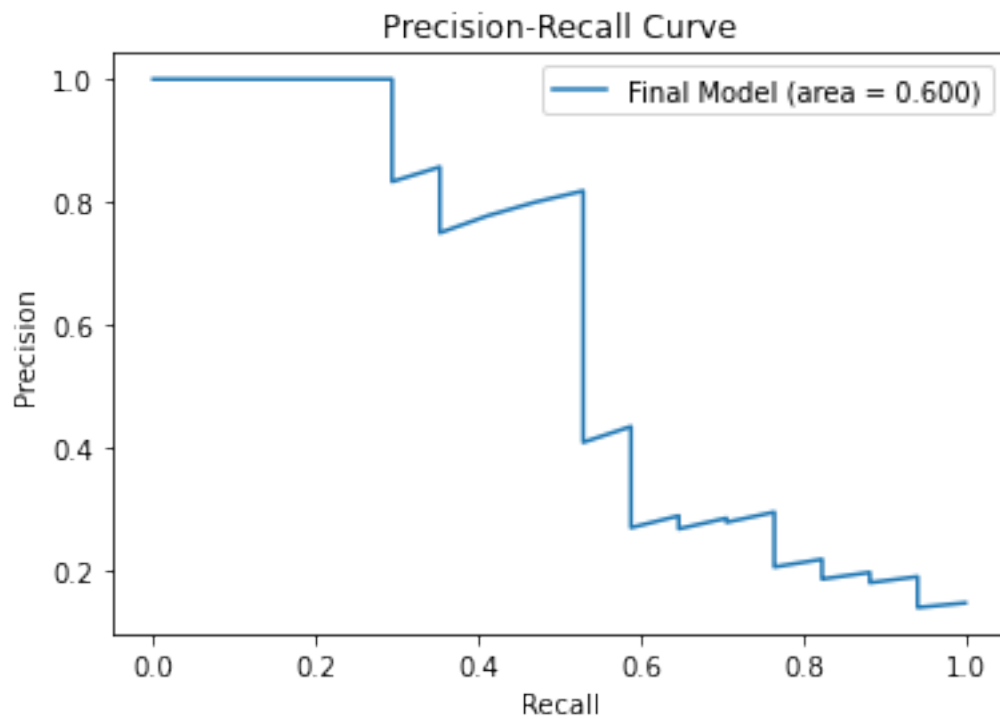
```
[ ]: # Plot ROC curve
plt.plot([0, 1], [0, 1], 'k--')
plt.plot(fpr, tpr, label='Final Model (area = {:.3f})'.format(model_auc))
plt.xlabel('False positive rate')
plt.ylabel('True positive rate')
plt.legend(loc='lower right')
plt.title('ROC Curve')
plt.show()
```



7.5 Precision-Recall Curve

```
[ ]: # Get precision-recall variables
precision, recall, thresholds = precision_recall_curve(test_y, test_pred)
model_auprc = auc(recall, precision)
```

```
[ ]: # Plot precision-recall curve
plt.plot(recall, precision, label='Final Model (area = {:.3f})'.
    ↳format(model_auprc))
plt.xlabel('Recall')
plt.ylabel('Precision')
plt.legend(loc='upper right')
plt.title('Precision-Recall Curve')
plt.show()
```

7.6 F1-Score

```
[ ]: # Get F1 score
f1 = f1_score(test_y, np.rint(test_pred))
print("The F1-score was %5.3f" % (f1))
```

The F1-score was 0.476